

# Aprendizado de Máquina e Reconhecimento de Padrões



## K-Nearest Neighbors (KNN)

Prof. Dr. Samuel Martins (Samuka)

[samuel.martins@ifsp.edu.br](mailto:samuel.martins@ifsp.edu.br)

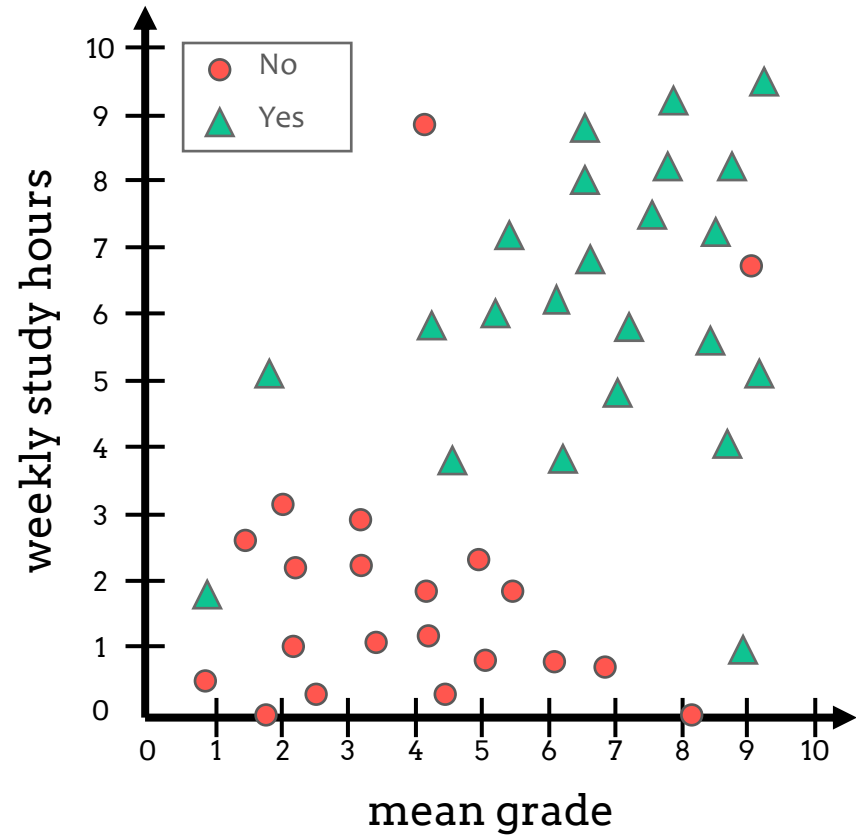


# KNN


One of the simplest (**instance-based**) classification algorithms.

It works for binary or multiclass problems.

Training Set		
mean grade	weekly study hours	approved at a university?
0.9	0.5	No
2.2	2	No
9.00	7.2	Yes
6.5	8.00	Yes
...	...	...

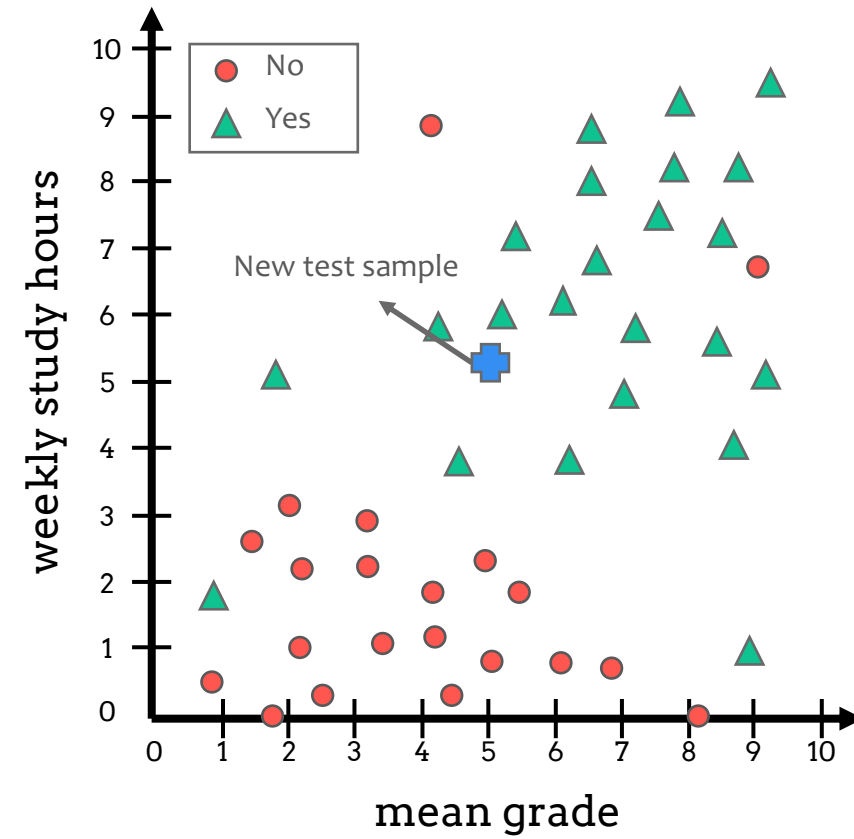


# No Training



 KNN **does not** learn any model.


# Classification



# Classification

hyperparameter

- **Step 1:** Choose the number **K** of neighbors;




**k = 3**

# Classification

hyperparameter

- **Step 1:** Choose the number **K of neighbors**;
- **Step 2:** Take the K nearest neighbors of the **new instance**, according to a given **distance measure** (e.g., Euclidean);

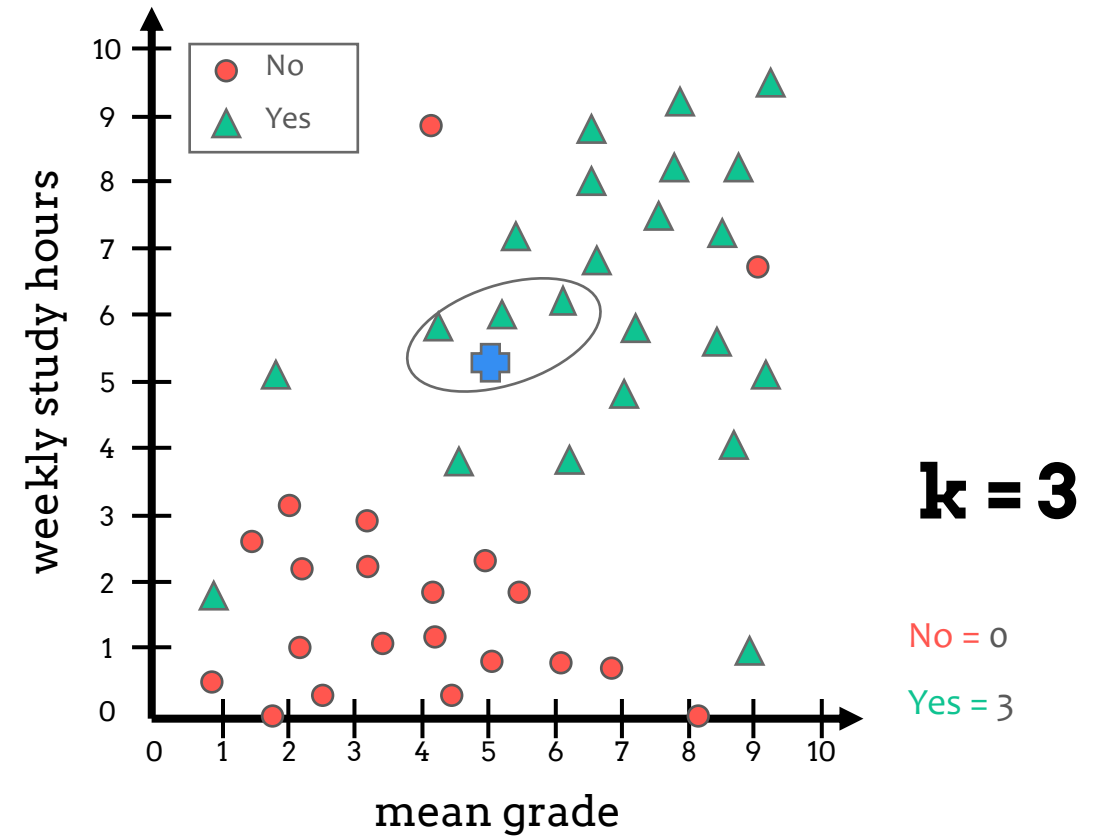


**k = 3**

# Classification

hyperparameter


- **Step 1:** Choose the number **K of neighbors**;
- **Step 2:** Take the K nearest neighbors of the **new instance**, according to a given **distance measure** (e.g., Euclidean);
- **Step 3:** Among these K neighbors, **count** the number of training instances in each class/category.



# Classification

hyperparameter

- **Step 1:** Choose the number **K of neighbors**;
- **Step 2:** Take the K nearest neighbors of the **new instance**, according to a given **distance measure** (e.g., Euclidean);
- **Step 3:** Among these K neighbors, **count** the number of training instances in each class/category.
- **Step 4:** Assign the **new test instance** to the **most frequent class** (majority voting).



Yes  
**k = 3**

No = 0


Yes = 3



# Classification

hyperparameter

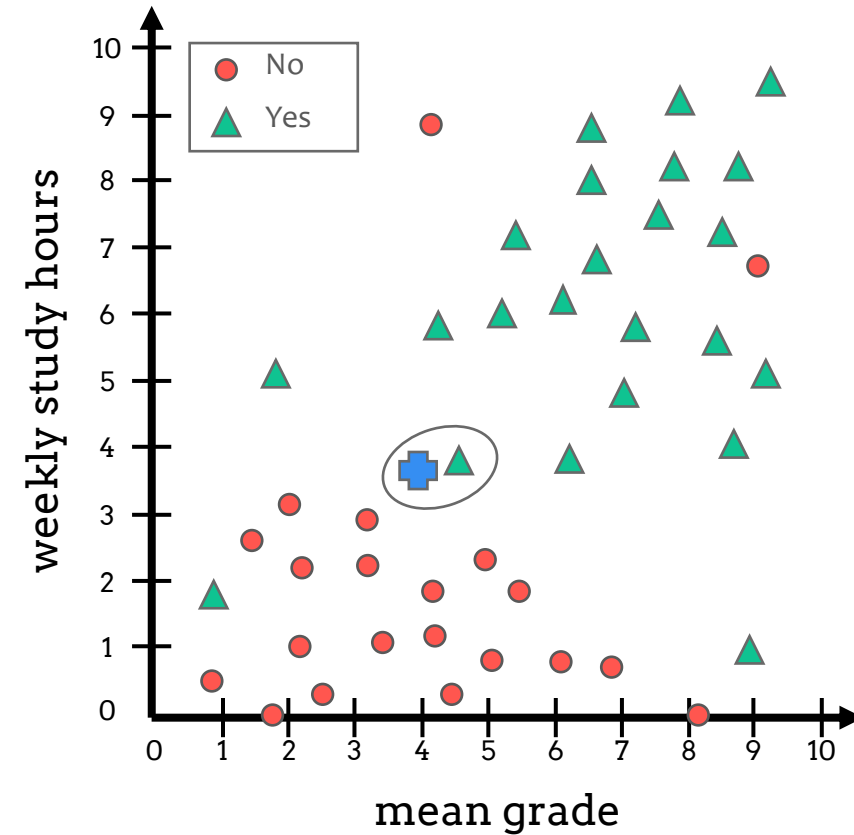
- **Step 1:** Choose the number **K of neighbors**;
- **Step 2:** Take the K nearest neighbors of the **new instance**, according to a given **distance measure** (e.g., Euclidean);
- **Step 3:** Among these K neighbors, **count** the number of training instances in each class/category.
- **Step 4:** Assign the **new test instance** to the **most frequent class** (majority voting).



# Classification

hyperparameter

- **Step 1:** Choose the number **K of neighbors**;
- **Step 2:** Take the K nearest neighbors of the **new instance**, according to a given **distance measure** (e.g., Euclidean);
- **Step 3:** Among these K neighbors, **count** the number of training instances in each class/category.
- **Step 4:** Assign the **new test instance** to the **most frequent class** (majority voting).



Yes

**k = 1**


No = 0

Yes = 1

# Classification

hyperparameter

- **Step 1:** Choose the number **K of neighbors**;
- **Step 2:** Take the K nearest neighbors of the **new instance**, according to a given **distance measure** (e.g., Euclidean);
- **Step 3:** Among these K neighbors, **count** the number of training instances in each class/category.
- **Step 4:** Assign the **new test instance** to the **most frequent class** (majority voting).



No  
**k = 3**

No = 2

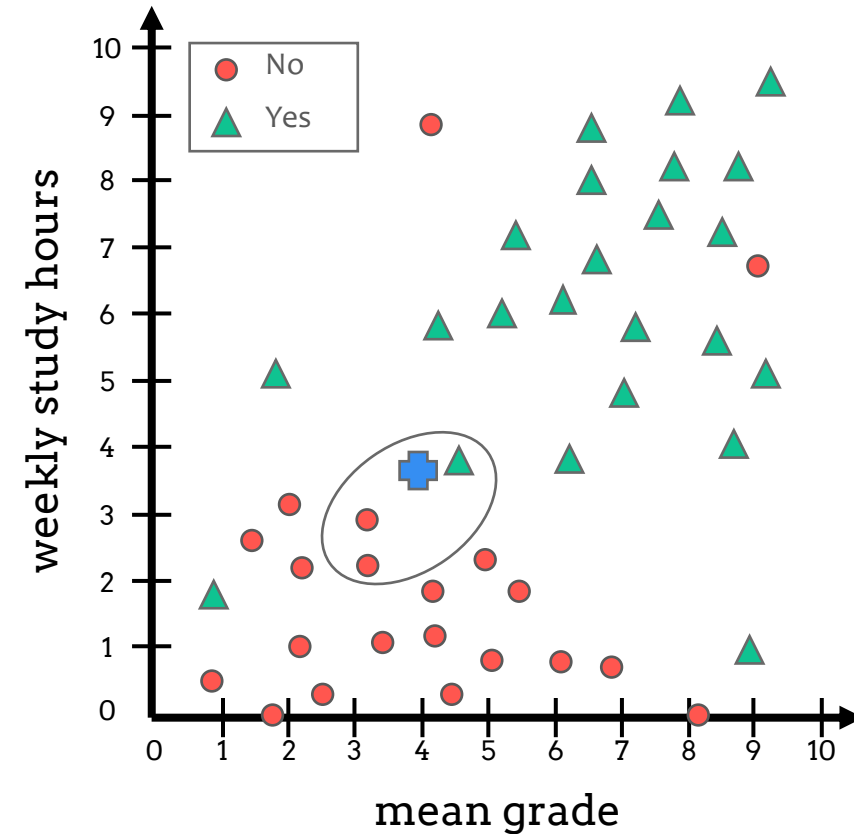
Yes = 1

# Classification

hyperparameter

- **Step 1:** Choose the number **K of neighbors**;
- **Step 2:** Take the K nearest neighbors of the **new instance**, according to a given **distance measure** (e.g., Euclidean);
- **Step 3:** Among these K neighbors, **count** the number of training instances in each class/category.
- **Step 4:** Assign the **new test instance** to the **most frequent class** (majority voting).

`sklearn.neighbors.KNeighborsClassifier`



**No**  
**k = 3**

No = 2

Yes = 1


# Common Distance Measures

---

# Common Distance Measures

---

Euclidean




$$d(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$


- Common distance measure;
- Suitable for low-dimensional data;

# Common Distance Measures

Euclidean



Manhattan



$$d(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$


- Common distance measure;
- Suitable for low-dimensional data;

$$d(a, b) = \sum_{i=1}^n |b_i - a_i|$$


- Work quite well when your data has discrete and/or binary attributes;
- Work ok for high-dimensional data;
- Less intuitive than Euclidean distance;
- In general, give a higher distance value than Euclidean distance;

# Common Distance Measures


Euclidean



Manhattan



Chebyshev



$$d(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$

- Common distance measure;
- Suitable for low-dimensional data;

$$d(a, b) = \sum_{i=1}^n |b_i - a_i|$$

- Work quite well when your data has discrete and/or binary attributes;
- Work ok for high-dimensional data;
- Less intuitive than Euclidean distance;
- In general, give a higher distance value than Euclidean distance;


$$d(a, b) = \max_i (|b_i - a_i|)$$

- It can be used to extract the minimum number of moves needed to get from one square to another;
- Used in very specific use-cases, such as warehouse logistics;




# Common Distance Measures


Euclidean




Manhattan



Chebyshev



Minkowski



$$d(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$

- Common distance measure;
- Suitable for low-dimensional data;

$$d(a, b) = \sum_{i=1}^n |b_i - a_i|$$

- Work quite well when your data has discrete and/or binary attributes;
- Work ok for high-dimensional data;
- Less intuitive than Euclidean distance;
- In general, give a higher distance value than Euclidean distance;

$$d(a, b) = \max_i (|b_i - a_i|)$$


- It can be used to extract the minimum number of moves needed to get from one square to another;
- Used in very specific use-cases, such as warehouse logistics;

$$d(a, b) = \left( \sum_{i=1}^n |b_i - a_i|^p \right)^{\frac{1}{p}}$$

- Metric in a normed vector space;
- The upside to  $p$  is the possibility to iterate over it and find the distance measure that works best for your use case.
- $p=1 \rightarrow$  Manhattan distance
- $p=2 \rightarrow$  Euclidean distance
- $p=\infty \rightarrow$  Chebyshev distance

# Common Distance Measures


Euclidean



$$d(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$


- Common distance measure;
- Suitable for low-dimensional data;

Manhattan



$$d(a, b) = \sum_{i=1}^n |b_i - a_i|$$


- Work quite well when your data has discrete and/or binary attributes;
- Work ok for high-dimensional data;
- Less intuitive than Euclidean distance;
- In general, give a higher distance value than Euclidean distance;



$$d(a, b) = \max_i (|b_i - a_i|)$$

- It can be used to extract the minimum number of moves needed to get from one square to another;
- Used in very specific use-cases, such as warehouse logistics;

Minkowski




$$d(a, b) = \left( \sum_{i=1}^n |b_i - a_i|^p \right)^{\frac{1}{p}}$$

- Metric in a normed vector space;
- The upside to p is the possibility to iterate over it and find the distance measure that works best for your use case.
- $p=1 \rightarrow$  Manhattan distance
- $p=2 \rightarrow$  Euclidean distance
- $p=\infty \rightarrow$  Chebyshev distance

# Common Distance Measures


Euclidean



$$d(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$


- Common distance measure;
- Suitable for low-dimensional data;

Manhattan



$$d(a, b) = \sum_{i=1}^n |b_i - a_i|$$


- Work quite well when your data has discrete and/or binary attributes;
- Work ok for high-dimensional data;
- Less intuitive than Euclidean distance;
- In general, give a higher distance value than Euclidean distance;



$$d(a, b) = \max_i (|b_i - a_i|)$$

- It can be used to extract the minimum number of moves needed to get from one square to another;
- Used in very specific use-cases, such as warehouse logistics;

Minkowski




$$d(a, b) = \left( \sum_{i=1}^n |b_i - a_i|^p \right)^{\frac{1}{p}}$$


- Metric in a normed vector space;
- The upside to p is the possibility to iterate over it and find the distance measure that works best for your use case.
- $p=1 \rightarrow$  Manhattan distance
- $p=2 \rightarrow$  Euclidean distance
- $p=\infty \rightarrow$  Chebyshev distance

# Common Distance Measures


Euclidean




Manhattan



Chebyshev



Minkowski



$$d(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$

- Common distance measure;
- Suitable for low-dimensional data;



Typically, these metrics require **feature scaling**.

$$d(a, b) = \sum_{i=1}^n |b_i - a_i|$$

- Work quite well when your data has discrete and/or binary attributes;
- Work ok for high-dimensional data;
- Less intuitive than Euclidean distance;
- In general, give a higher distance value than Euclidean distance;

$$d(a, b) = \max_i (|b_i - a_i|)$$


- It can be used to extract the minimum number of moves needed to get from one square to another;
- Used in very specific use-cases, such as warehouse logistics;

$$d(a, b) = \left( \sum_{i=1}^n |b_i - a_i|^p \right)^{\frac{1}{p}}$$


- Metric in a normed vector space;
- The upside to  $p$  is the possibility to iterate over it and find the distance measure that works best for your use case.
- $p=1 \rightarrow$  Manhattan distance
- $p=2 \rightarrow$  Euclidean distance
- $p=\infty \rightarrow$  Chebyshev distance

# KNN with and without Feature Scaling

Data without Feature Scaling



Data with Feature Scaling



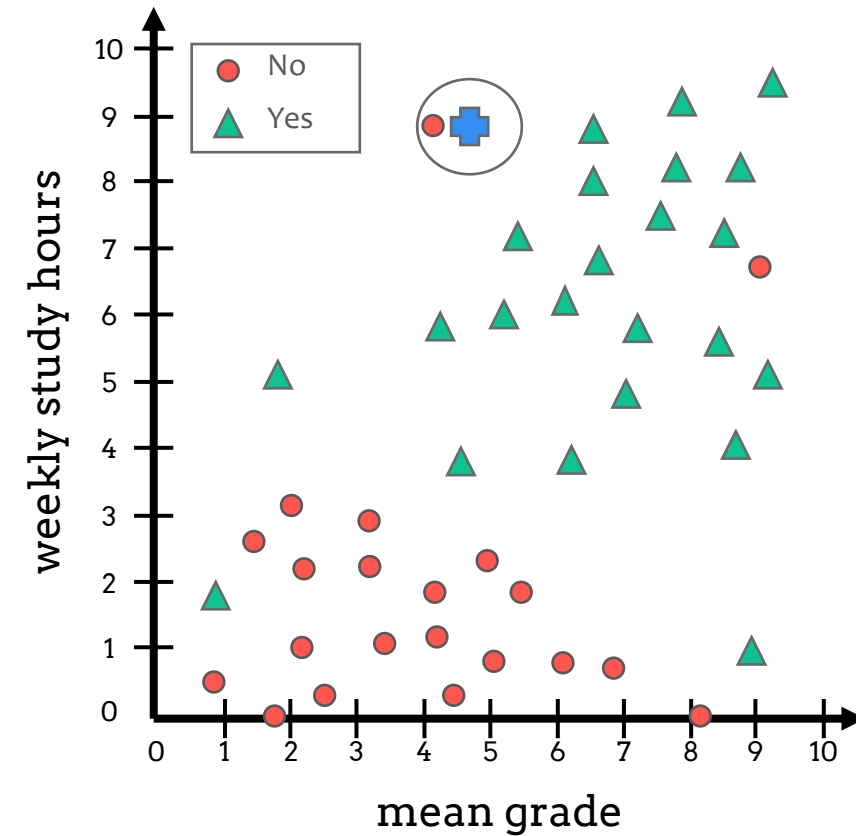
KNN works properly after  
**feature scaling** your data.

# Choosing the value of $K$

---


# Choosing the value of K

- If **K** is too small:
  - KNN is sensitive to *outliers*;
  - KNN overfits the training data:
    - Higher error on different sets;



# Choosing the value of K


- If **K is too small:**
  - KNN is sensitive to *outliers*;
  - KNN overfits the training data:
    - Higher error on different sets;
- If **K is too large:**
  - Neighbors can include samples from other classes;
  - KNN performs poorly on both train and validation set;






# Choosing the value of K

- If **K is too small**:
  - KNN is sensitive to *outliers*;
  - KNN overfits the training data:
    - Higher error on different sets;
- If **K is too large**:
  - Neighbors can include samples from other classes;
  - KNN performs poorly on both train and validation set;
- **K** is a **hyperparameter** so we can **optimize it** for our problem;
  - For example, by **Cross-Validation Grid Search**



# Choosing the value of K

- If **K is too small**:
  - KNN is sensitive to *outliers*;
  - KNN overfits the training data:
    - Higher error on different sets;
- If **K is too large**:
  - Neighbors can include samples from other classes;
  - KNN performs poorly on both train and validation set;
- **K is a hyperparameter** so we can **optimize it** for our problem;
  - For example, by **Cross-Validation Grid Search**
  - By plotting an **elbow curve**



# KNN: Pros and Cons

---

## Pros

1. Extremely easy to implement it;
2. It does not require training;
3. By **not** requiring **training** before making estimation/classifications, **new training samples** can be added without any problems (no models' retraining);
4. There is **only a single hyperparameter** required by KNN
  - Number of neighbors K
  - If we consider other distances, we can have more required hyperparameters (Minkowski, ...);

# KNN: Pros and Cons

---

## Pros


1. Extremely easy to implement it;
2. It does not require training;
3. By **not** requiring **training** before making estimation/classifications, **new training samples** can be added without any problems (no models' retraining);
4. There is **only a single hyperparameter** required by KNN
  - Number of neighbors K
  - If we consider other distances, we can have more required hyperparameters (Minkowski, ...);

## Cons

1. It does not work well for **high dimensionality data (the curse of dimensionality)**
2. The **prediction time** can be **high** if the size of the training set is large;


# Variant: KNN Regression

---



# Variant: KNN Regression


The **outcome** (label / dependent variable) of a **new test sample** is computed based on the **mean of the outcomes/labels** of its  $K$  nearest neighbors.



**$k = 3$**


# Variant: KNN Regression

The **outcome** (label / dependent variable) of a **new test sample** is computed based on the **mean of the outcomes/labels** of its  $K$  nearest neighbors.



# Variant: KNN Regression

The **outcome (label / dependent variable)** of a **new test sample** is computed based on the **mean of the outcomes/labels** of its  $K$  nearest neighbors.




**$k = 3$**



# Variant: KNN Regression

The **outcome** (label / dependent variable) of a **new test sample** is computed based on the **mean of the outcomes/labels** of its  $K$  nearest neighbors.



**$k = 3$**

```
sklearn.neighbors.KNeighborsRegressor
```

# Aprendizado de Máquina e Reconhecimento de Padrões



## K-Nearest Neighbors (KNN)

Prof. Dr. Samuel Martins (Samuka)

[samuel.martins@ifsp.edu.br](mailto:samuel.martins@ifsp.edu.br)

